

Introduction

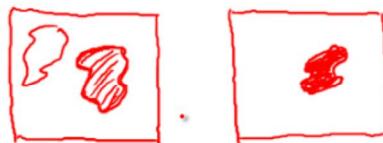
- Task

Divide an image into constituent “meaningful” regions

color
texture
motion

- Challenges

Image data ambiguity, noise, integration of high level information, lack of ground truth



- Approaches based on

intensity discontinuity, intensity similarity, morphology



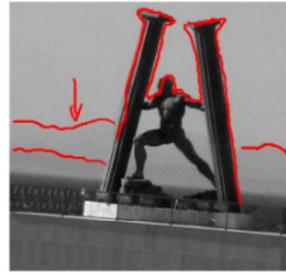
Agenda

- Methods based on intensity discontinuity
 - Edge-based segmentation
- Methods based on intensity similarity
 - Thresholding
 - Region growing
 - Region splitting and merging
- Watersheds and K-means algorithms
- Advanced methods

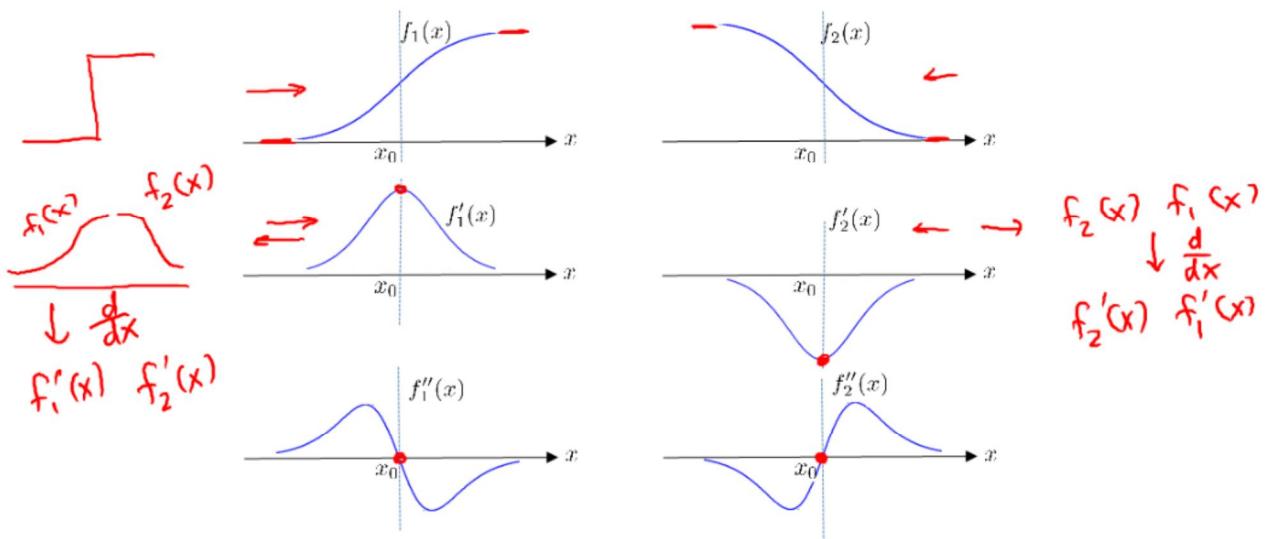


Edge Detection

- • Gradient-Based Approaches
 - ✓ Sobel
 - ✓ Prewitt
 - ✓ Roberts
- • Laplacian-Based Approaches
 - Marr-Hildreth algorithm
- Canny Algorithm



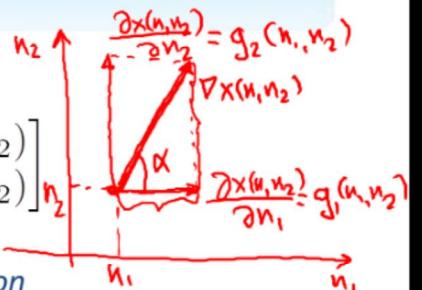
1st and 2nd Order Derivatives



The Image Gradient

Gradient operator

$$g_x(n_1, n_2) = \nabla x(n_1, n_2) = \begin{bmatrix} \frac{\partial x(n_1, n_2)}{\partial n_1} \\ \frac{\partial x(n_1, n_2)}{\partial n_2} \end{bmatrix} = \begin{bmatrix} g_1(n_1, n_2) \\ g_2(n_1, n_2) \end{bmatrix}$$



The Gradient magnitude is commonly used for edge detection

$$M(n_1, n_2) = |g_x(n_1, n_2)| = \left(\left(\frac{\partial x(n_1, n_2)}{\partial n_1} \right)^2 + \left(\frac{\partial x(n_1, n_2)}{\partial n_2} \right)^2 \right)^{\frac{1}{2}}$$

The direction of the gradient vector is given by

$$\alpha(n_1, n_2) = \tan^{-1} \left[\frac{g_2(n_1, n_2)}{g_1(n_1, n_2)} \right]$$



Gradient Evaluation

- How do we find the gradient?
- To approximate partial derivatives we use neighborhood differences.

$$\frac{\partial x(n_1, n_2)}{\partial n_1} = g_1(n_1, n_2) = x(n_1, n_2) * h_1(n_1, n_2) \quad ?$$

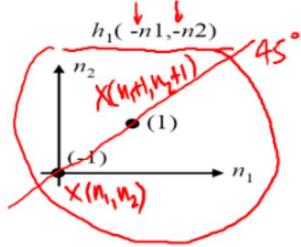
$$g_2(n_1, n_2) = x(n_1, n_2) * h_2(n_1, n_2) \quad ?$$

$$M(n_1, n_2) = |g(n_1, n_2)| = \sqrt{g_1(n_1, n_2)^2 + g_2(n_1, n_2)^2}$$

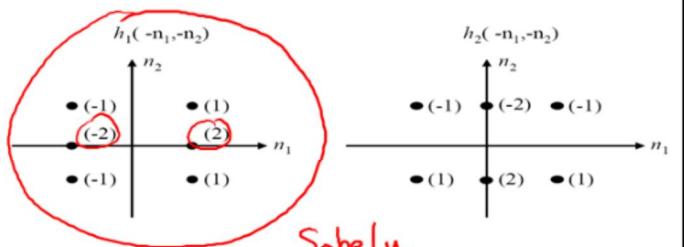


Gradient Kernels

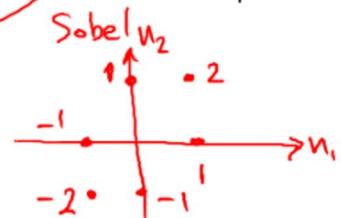
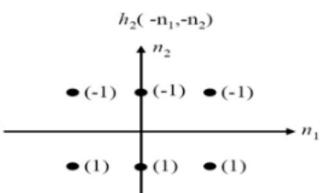
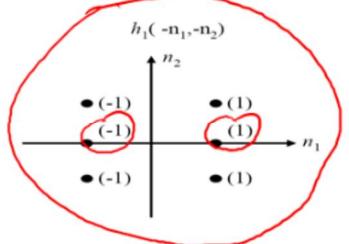
Roberts



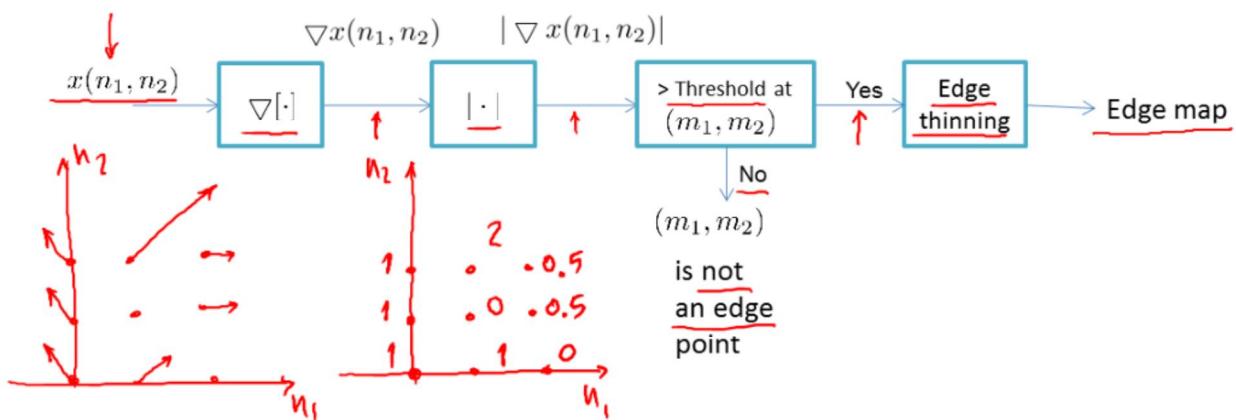
Sobel



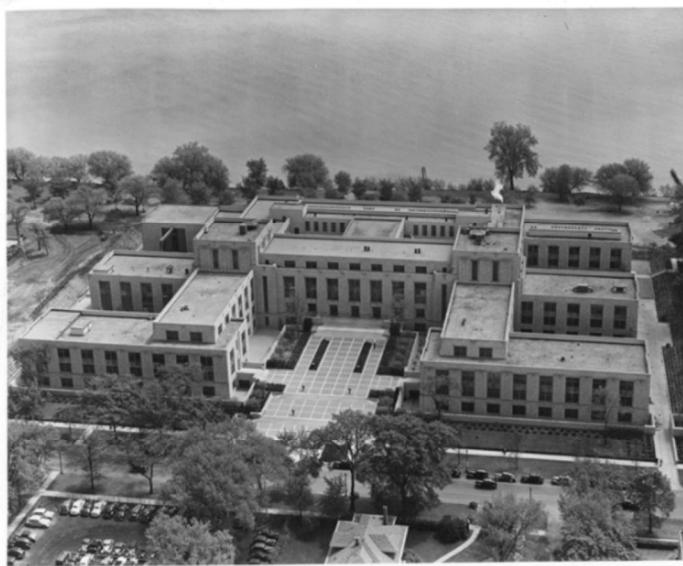
Prewitt



System for Edge Detection



Example



Technological Institute
Northwestern University
1942

Example



$\ast\ast h_1(u_1, u_2)$
=
Sobel

Technological Institute
Northwestern University
1942

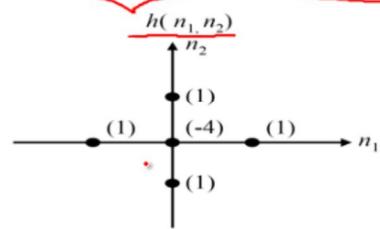
Example



Technological Institute
Northwestern University
1942

Calculation of the Laplacian

$$\nabla^2 x(n_1, n_2) = \frac{\partial^2 x(n_1, n_2)}{\partial n_1^2} + \frac{\partial^2 x(n_1, n_2)}{\partial n_2^2}$$
$$L(n_1, n_2) = x(n_1, n_2) * h_1(n_1, n_2) * h'_1(n_1, n_2) +$$
$$x(n_1, n_2) * h_2(n_1, n_2) * h'_2(n_1, n_2) +$$
$$L(n_1, n_2) = x(n_1, n_2) * (h_1(n_1, n_2) * h'_1(n_1, n_2) + h_2(n_1, n_2) * h'_2(n_1, n_2))$$
$$L(n_1, n_2) = x(n_1, n_2) * h(n_1, n_2)$$



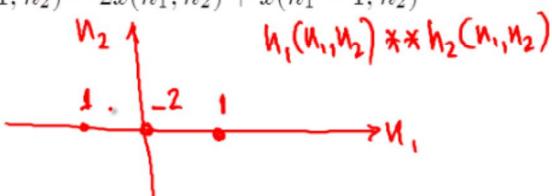
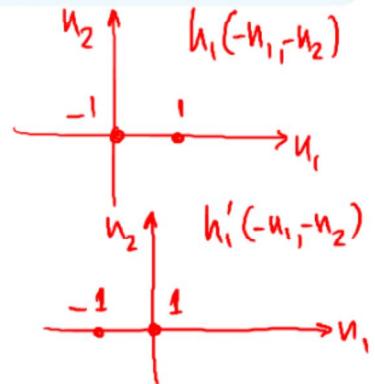
Calculation of the Laplacian

- Horizontal second derivative

$$\frac{\partial x(n_1, n_2)}{\partial n_1} = g_1(n_1, n_2) = x(n_1 + 1, n_2) - x(n_1, n_2)$$

$$\frac{\partial^2 x(n_1, n_2)}{\partial n_1^2} = \frac{\partial}{\partial n_1} \left(\frac{\partial x(n_1, n_2)}{\partial n_1} \right) = g_1(n_1, n_2) - g_1(n_1 - 1, n_2)$$

$$= x(n_1 + 1, n_2) - 2x(n_1, n_2) + x(n_1 - 1, n_2)$$



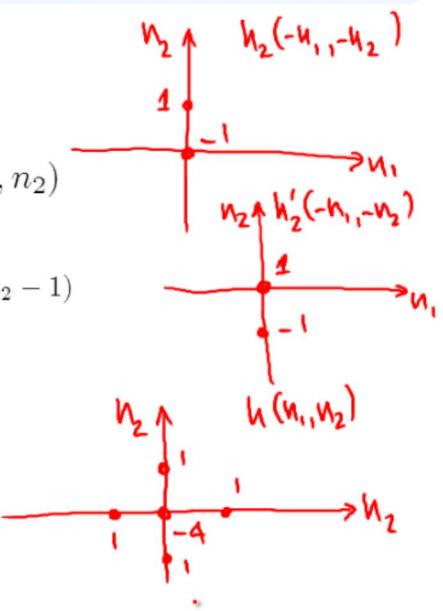
Calculation of the Laplacian

- Vertical second derivative

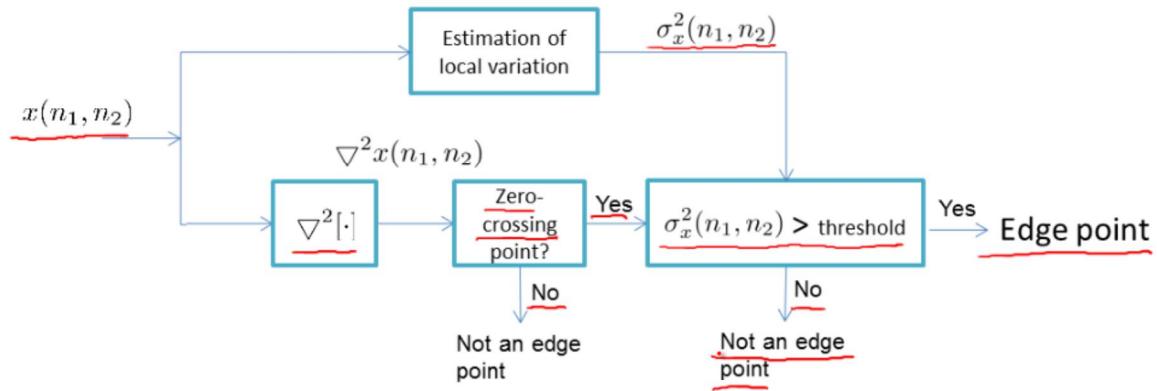
$$\frac{\partial x(n_1, n_2)}{\partial n_2} = g_2(n_1, n_2) = x(n_1, n_2 + 1) - x(n_1, n_2)$$

$$\frac{\partial^2 x(n_1, n_2)}{\partial n_2^2} = \frac{\partial}{\partial n_2} \left(\frac{\partial x(n_1, n_2)}{\partial n_2} \right) = g_2(n_1, n_2) - g_2(n_1, n_2 - 1)$$

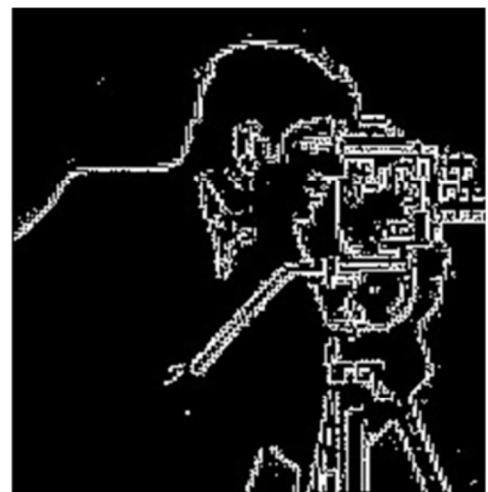
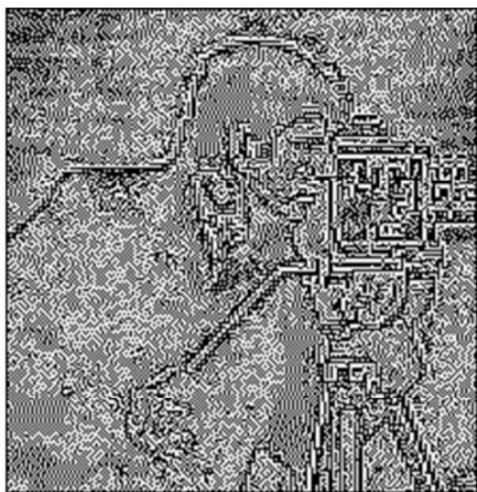
$$= x(n_1, n_2 + 1) - 2x(n_1, n_2) + x(n_1, n_2 - 1)$$



Laplacian-Based Edge Detection System



Example



Laplacian of Gaussian (LOG) Edge Detector

- Intensity changes occur at different scales
- Image is band-limited before edge detection
- Gaussian LPF

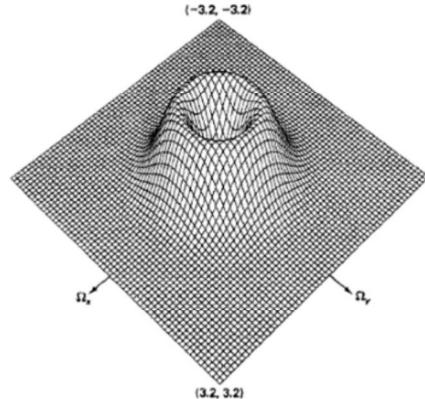
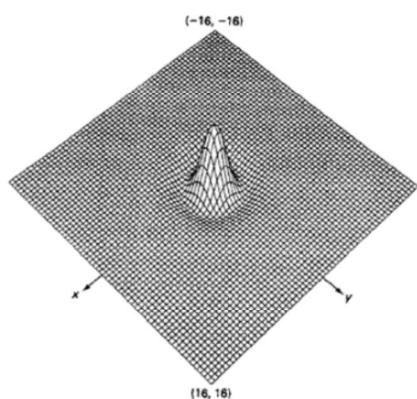
$$h(x, y) = e^{-\frac{(x^2+y^2)}{2\pi\sigma^2}}$$

- Taking the Laplacian

$$\nabla^2 (f(x, y) * h(x, y)) = f(x, y) * \underbrace{[\nabla^2 h(x, y)]}_{\text{LOG}}$$

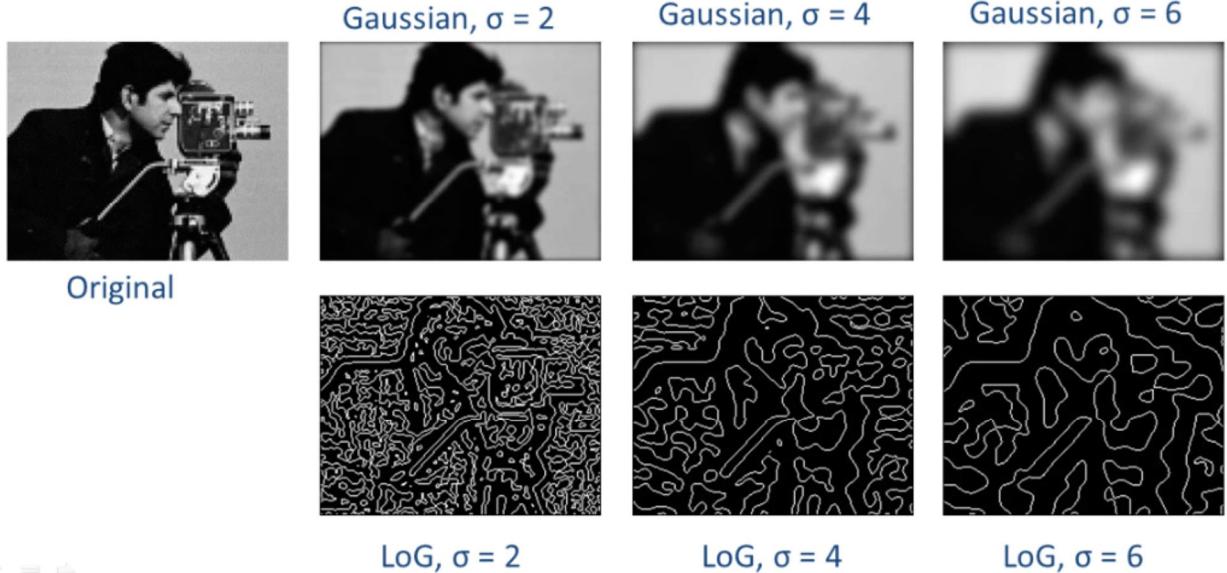
Laplacian of Gaussian (LOG) Edge Detector

$$\nabla^2 h(x, y) = \frac{e^{-\frac{(x^2+y^2)}{2\pi\sigma^2}}}{(\pi\sigma^2)^2} (x^2 + y^2 - 2\pi\sigma^2) \quad -FT\{\nabla^2 h(x, y)\} = 2\pi^2\sigma^2 \cdot (\Omega_x^2 + \Omega_y^2) \cdot e^{-\pi\sigma^2 \frac{(\Omega_x^2+\Omega_y^2)}{2}}$$



J. S. Lim, *Two-Dimensional Signal and Image Processing*, Prentice Hall, 1990.

Examples



The Canny Edge Detector

- ✓ Smooth the image with a Gaussian filter (σ)
- ✓ Compute the gradient magnitude and angle images
- Apply non-maximal suppression to the gradient magnitude image
- Use double thresholding (T_H, T_L)

→ J. Canny, "A computational approach to edge detection," *IEEE Trans. Pattern Analysis and Machine Intelligence*, No. 6, pp. 679–698, 1986.

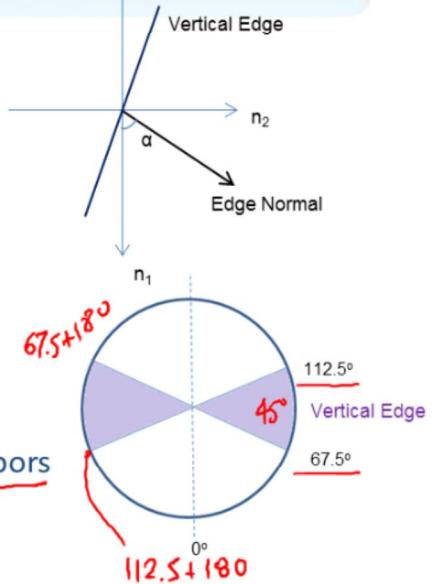
The Canny Edge Detector

Non-maximal suppression

Let d_1, d_2, d_3, d_4 denote the four basic edge directions:
horizontal, -45° , $+45^\circ$, vertical.

At every point (n_1, n_2) in $\alpha(n_1, n_2)$:

1. Find the direction d_k that is closest to $\alpha(n_1, n_2)$;
2. If the value of $M(n_1, n_2)$ is less than at least one of its neighbors along d_k , ignore/suppress it; otherwise, keep it.



The Canny Edge Detector

Double thresholding ($T_H > T_L$)

Pixels with values greater than T_H are “strong” edge pixels

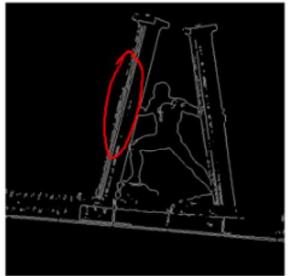
Pixels with values smaller than T_L are “weak” edge pixels

Edge Linking

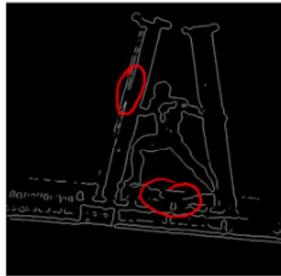
Locate a pixel p belonging to a “strong” edge;

Find “weak” pixels that are connected to p ;
append these “weak” pixels to p .

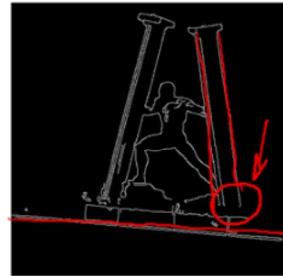
Example



Sobel



LOG



Canny

The Hough Transform

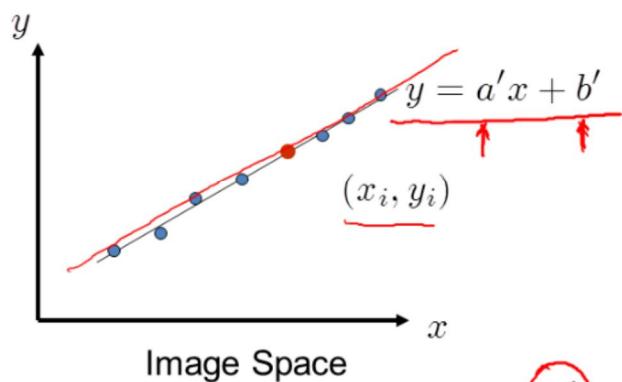
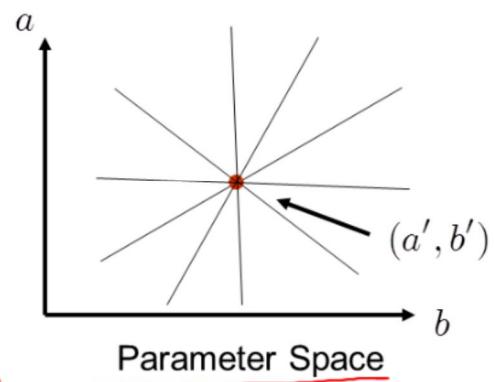


Image Space

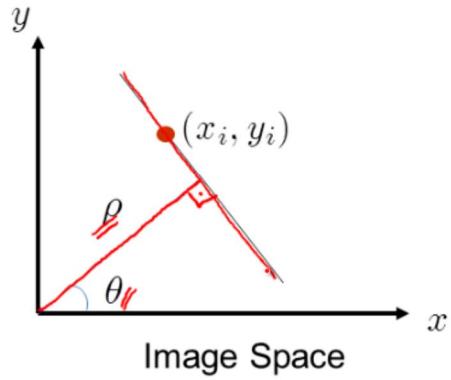


Parameter Space

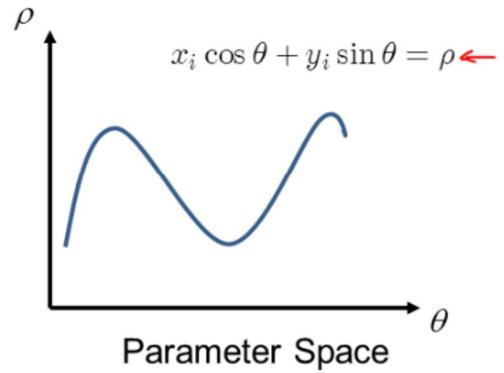
$$a = -\frac{1}{x}b + \frac{y}{x}$$

slope *a-intercept*

The Hough Transform



$$x \cos \theta + y \sin \theta = \rho$$



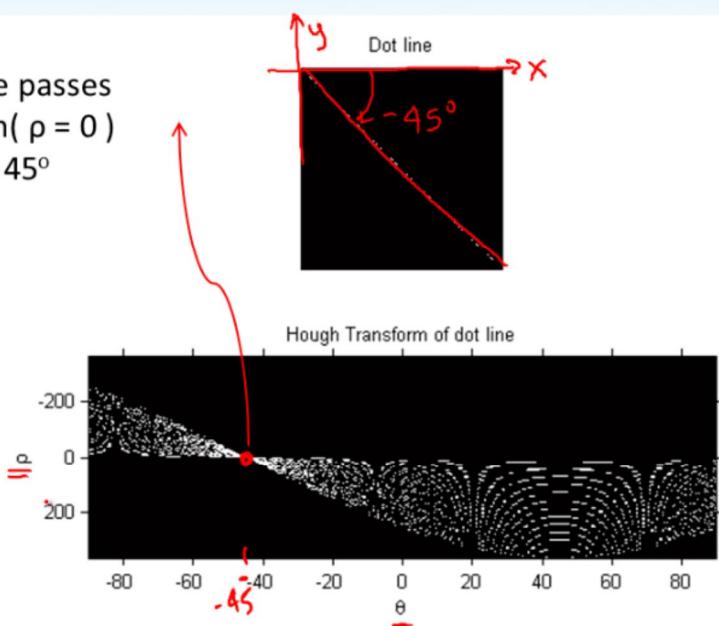
Parameter Space

$$x_i \cos \theta + y_i \sin \theta = \rho \leftarrow$$



The Hough Transform

The dominant line passes through the origin ($\rho = 0$) and oriented at -45°



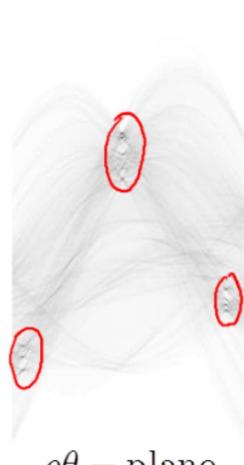
The Hough Transform

The Procedure:

- Obtain a binary edge image using any techniques (such as the ones discussed);
- Quantify the $\rho\theta$ -plane;
- Examine the number of intersects at each cell in the $\rho\theta$ -plane;
- Bridge the gap based on continuity*

**The gap in a line associated with a given cell is bridged if the length of the gap is less than a threshold.*

Example



Thresholding

- Two class case

$$x(n_1, n_2) = \begin{cases} \text{object} & x(n_1, n_2) \geq T \\ \text{background} & x(n_1, n_2) < T \end{cases}$$

- Approach can be easily applied when multiple thresholds are available
- The key of the thresholding process is the choice of the threshold value(s)

Example



Threshold $T = 0$

$$x(n_1, n_2) = \begin{cases} 1 & x(n_1, n_2) \geq T \\ 0 & x(n_1, n_2) < T \end{cases}$$

Example



Threshold T = 10

$$x(n_1, n_2) = \begin{cases} 1 & x(n_1, n_2) \geq T \\ 0 & x(n_1, n_2) < T \end{cases}$$

Example



Threshold T = 25

$$x(n_1, n_2) = \begin{cases} 1 & x(n_1, n_2) \geq T \\ 0 & x(n_1, n_2) < T \end{cases}$$

Example



Threshold T = 50

$$x(n_1, n_2) = \begin{cases} 1 & x(n_1, n_2) \geq T \\ 0 & x(n_1, n_2) < T \end{cases}$$



Example

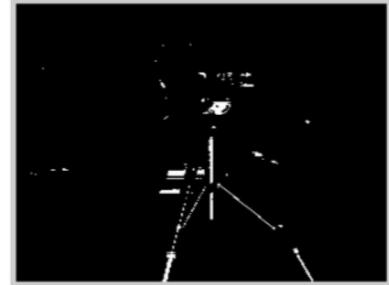


Threshold T = 150

$$x(n_1, n_2) = \begin{cases} 1 & x(n_1, n_2) \geq T \\ 0 & x(n_1, n_2) < T \end{cases}$$



Example



$$x(n_1, n_2) = \begin{cases} 1 & x(n_1, n_2) \geq T \\ 0 & x(n_1, n_2) < T \end{cases}$$

Otsu's Method

- ✓ Automatically finds the optimal global thresholding maximizing the between-class variance
- ✓ Only the histogram of the image is used

N. Otsu, "A threshold selection method from gray-level histograms", *Automatica*, 11(285-296):23–27, 1975.

Otsu's Method

- Compute the normalized histogram of the input image. Denote the components of the histogram by $p_i, i = 0, 1, 2, \dots, L - 1$
- Suppose a threshold is selected $k, 0 < k < L - 1$
- C_1 is the set of pixels with levels $[0, 1, 2, \dots, k]$
- C_2 is the set of pixels with levels $[k + 1, \dots, L - 1]$
- Obtain the value of the threshold which maximizes the between class variance
$$\rightarrow \sigma_B^2(k) = P_1(k)(m_1(k) - \underline{m}_G)^2 + P_2(k)(m_2(k) - \underline{m}_G)^2$$

Otsu's Method

- $P_1(k)$ is the probability of the set C_1 occurring
$$P_1(k) = \sum_{i=0}^k p_i, k = 0, 1, 2, \dots, L - 1$$
$$P_2(k) = \sum_{i=k+1}^{L-1} p_i = 1 - P_1(k), k = 0, 1, 2, \dots, L - 1$$
- $m_1(k)$ and $m_2(k)$ are the mean intensities of C_1 and C_2 , respectively
- The mean intensity up to level k is given by
$$m(k) = \sum_{i=0}^k ip_i, k = 0, 1, 2, \dots, L - 1$$
- Global intensity mean $m_G = \sum_{i=0}^{L-1} ip_i$

Otsu's Method

- Substituting the expressions from the previous slide we obtain

$$\sigma_B^2(k) = \frac{[m_G \tilde{P}_1(k) - \tilde{m}(k)]^2}{\tilde{P}_1(k)[1 - \tilde{P}_1(k)]}, \quad k = 0, 1, 2, \dots, L - 1$$

\uparrow

- Only two parameters have to be computed for all values of k
- The Otsu threshold \underline{k}^* is the value of k which maximizes the *between-class variance* above.

Otsu's Method Example



Region-Based Segmentation

- Basic Formulation: Let R represent the entire image region. Segmentation is the process of partitioning R into subregions R_1, R_2, \dots, R_n so that

$$\bigcup_{i=1}^n R_i = R$$

R_i is a connected region, $i = 1, 2, \dots, n$

$R_i \bigcap R_j = \emptyset$ for all i and $j, i \neq j$

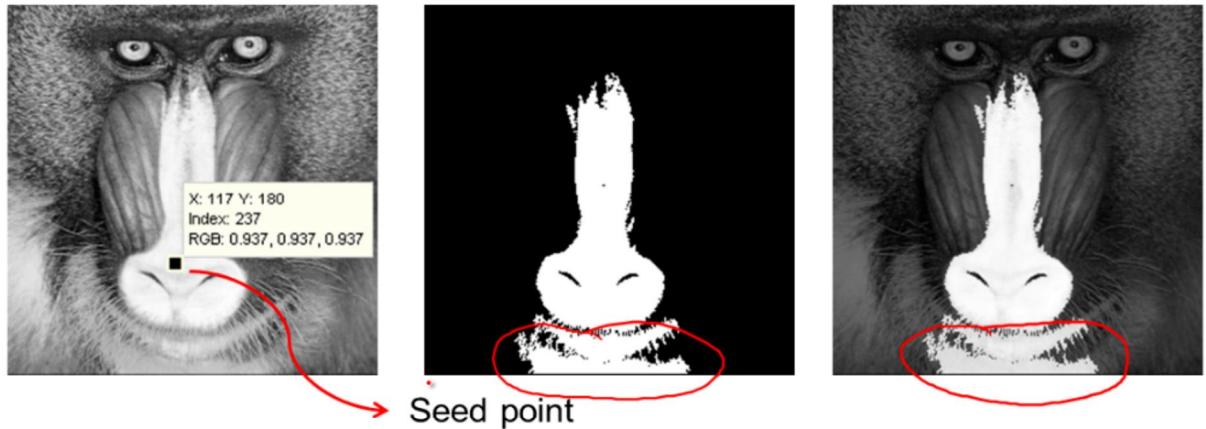
$P(R_i) = \text{TRUE}$ for $i = 1, 2, \dots, n$

$P(R_i \cup R_j) = \text{FALSE}$ for $i = 1, 2, \dots, n$

Region Growing

- A set of “seed” points is selected
 - Based on prior information
 - Based on properties computed at each pixel
- Regions grow from “seed” points by appending to each seed those neighboring pixels that have predefined properties similar to the seed
 - Similarities in color, moments, texture, etc.

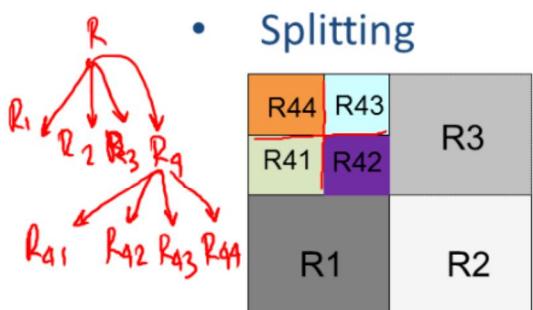
Region Growing Example



$$S(x(n_1, n_2), x(m_1, m_2)) = \exp(-|x(n_1, n_2) - x(m_1, m_2)|^2)$$

Region Splitting and Merging

- Splitting



\underline{R} represents the entire image
Select a predicate function \underline{P}
Divide the image into smaller quadrants if
 $\underline{P(R)} = \underline{\text{FALSE}}$
A minimum quad-region size is selected.

- Merging

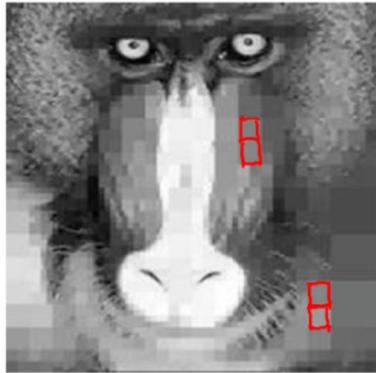
Two adjacent regions R_j and R_k are merged if

$$\underline{P(R_j \cup R_k)} = \underline{\text{TRUE}}$$

Region Splitting and Merging



Original Baboon Image



Region splitting



Region splitting and merging

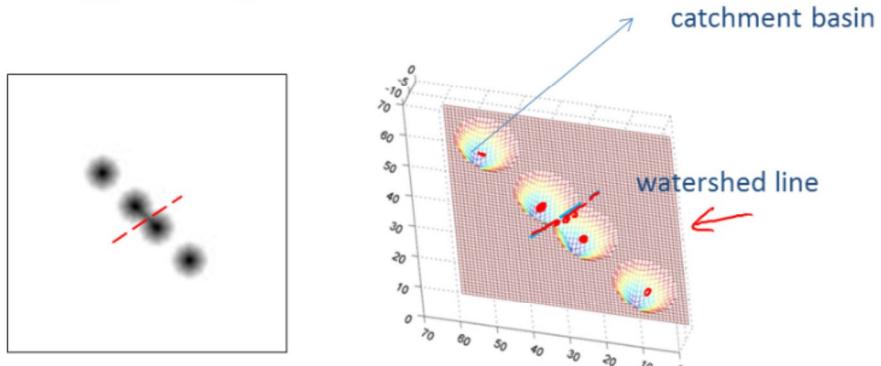
Region Splitting and Merging



(2x2)

Watersheds

- Watersheds segmentation envisions the grayscale image as a topological surface



Three types of points

- Points belonging to a **regional minimum**
- **Catchment basin** of a regional minimum
Points at which a drop of water will certainly fall to a single minimum
- Ridge lines / **Watershed lines**
Points at which a drop of water will be equally likely to fall to more than one minimum

Goal of the algorithm

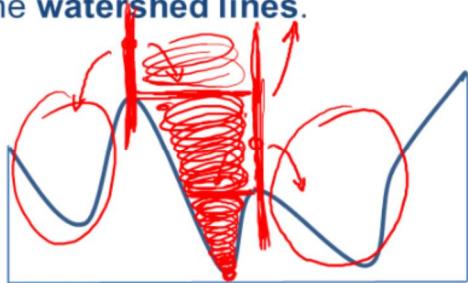
- Find the **watershed lines**

Watersheds

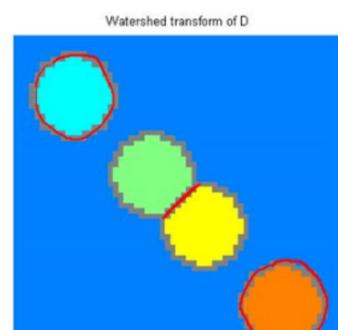
The concept:

Suppose that a hole is punched in each regional minimum and that the entire topography is flooded from below by letting water rise through the holes at a uniform rate.

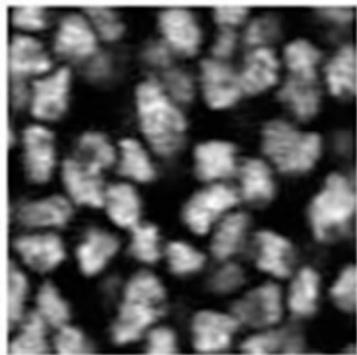
When rising water in distinct catchment basins is about to merge, a dam is built to prevent merging. These dam boundaries correspond to the **watershed lines**.



Example



Watersheds Example



Watersheds

A solution to the resulting over-segmentation is to use markers to specify the only allowed regional minima.

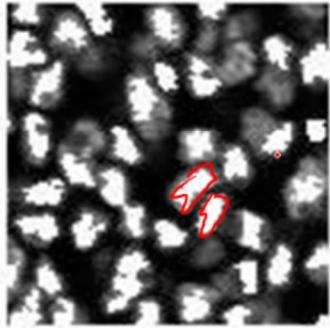
Internal markers:

- each one corresponds to one object
- surrounded by points of higher “altitude”
- points in region form a connected component
- all points in connected components have the same intensity

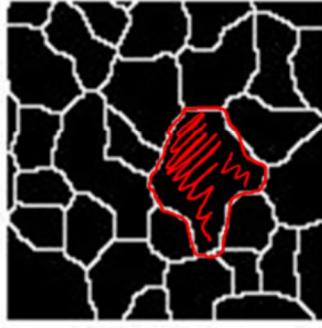
External markers:

- partition the image into regions where the regional minima are allowed to locate

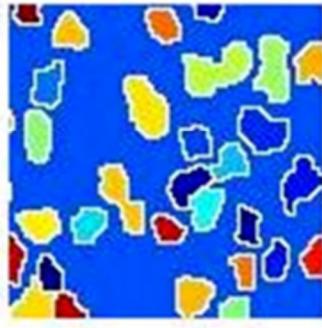
Watersheds Example



Internal markers overlaid
on the original image



External markers



Watersheds with markers



K-means Clustering

Clustering means grouping similar data together

From a set of data points or observations (all numerical), K-means attempts to classify all data points into K clusters

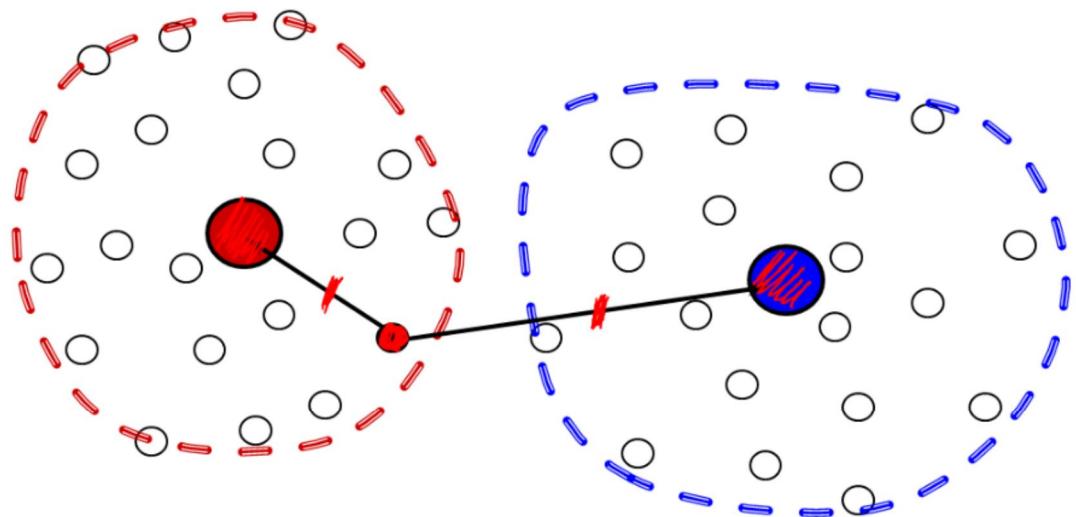
K clusters may represent K segments of an image

“K” stands for number of clusters, it is a user input to the algorithm

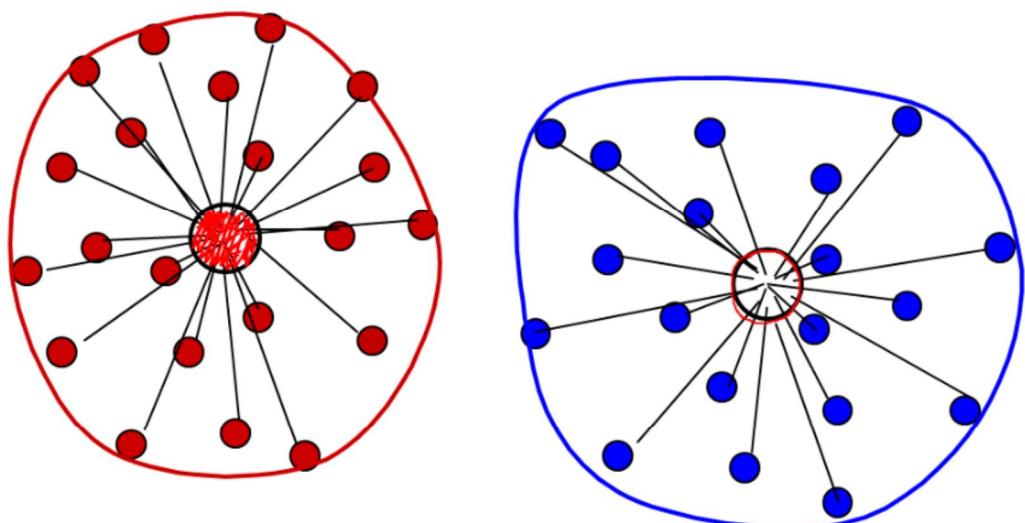
The algorithm is iterative in nature and not confined to image segmentation



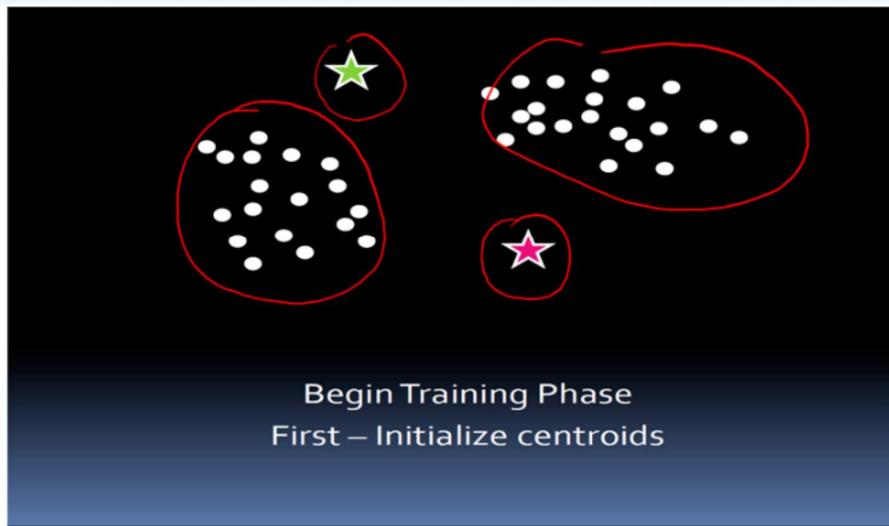
Sub-problem I: Clustering by distance to known centers



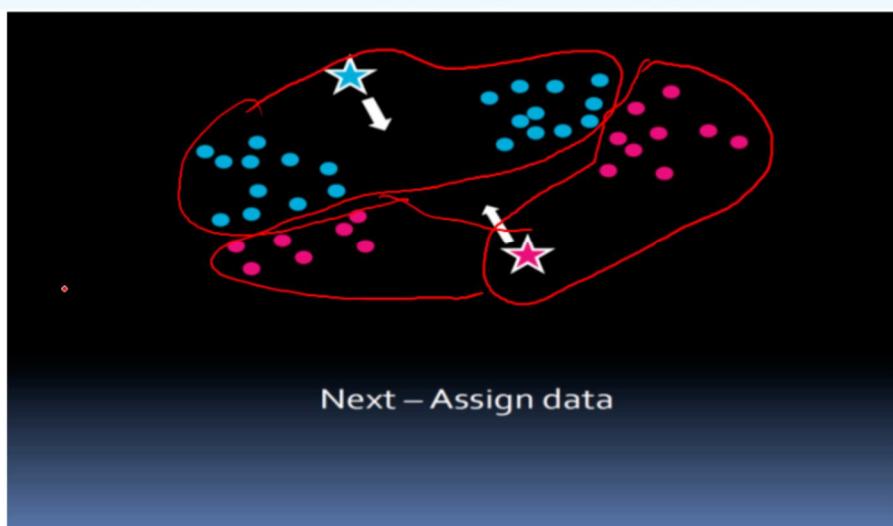
Sub-problem II: Finding new centers from known clustering



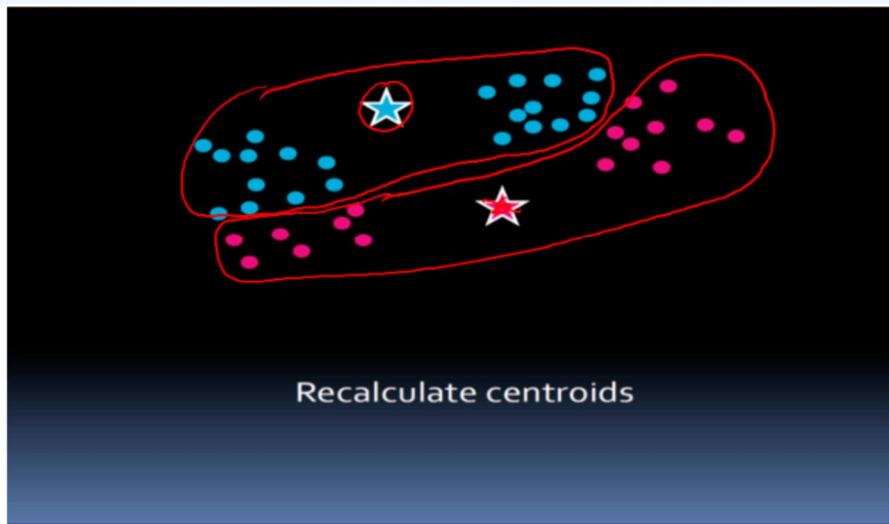
The K-means Clustering



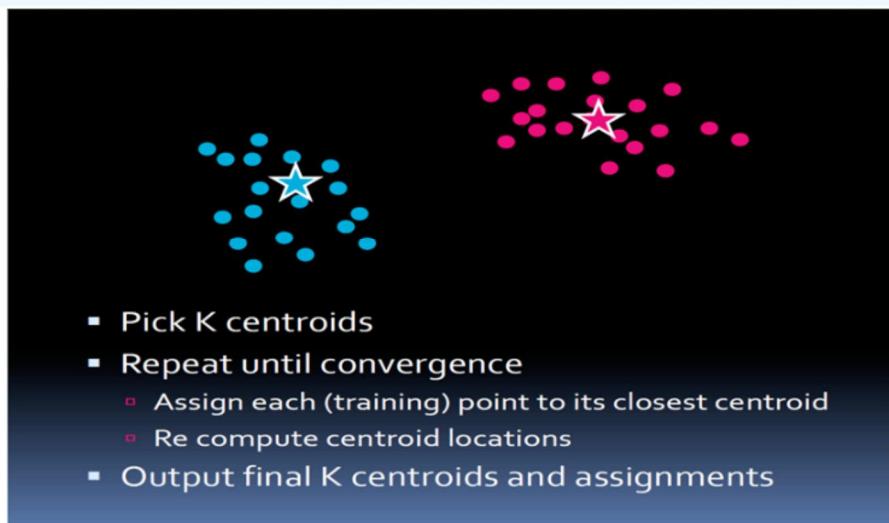
The K-means Clustering



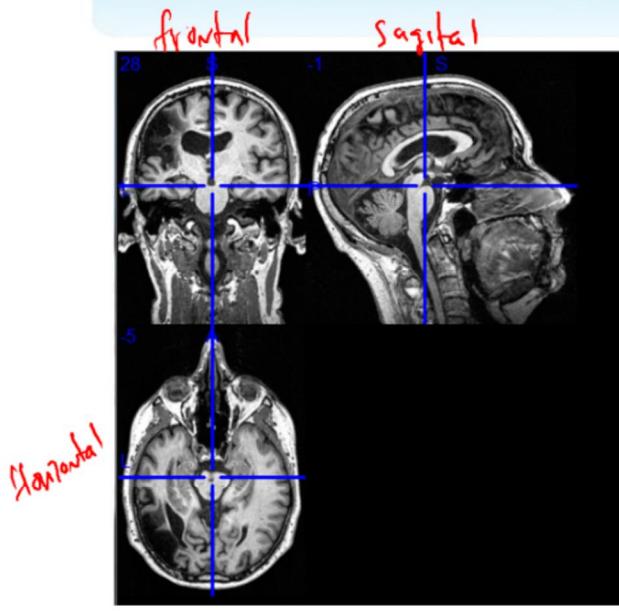
The K-means Clustering



The K-means Clustering



The K-means Clustering



$$v(x, y, z) = \begin{bmatrix} I(x, y, z) \\ x \\ y \\ z \end{bmatrix}$$

$$D(v_i, v_j) = \|v_i - v_j\|$$

Prior knowledge:

- Voxel with same intensity should belong to the same group;
- Voxel next to each other should belong to the same group;
- There are three groups: White matter, Grey matter and CSF

The K-means Clustering

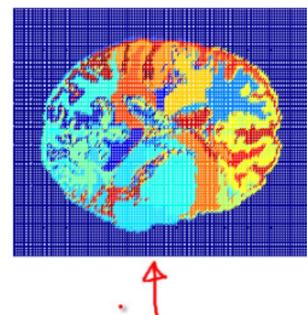
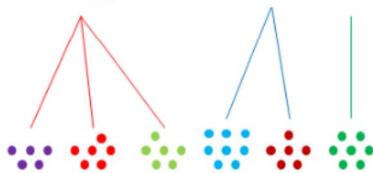
- The definition of vector and distance

$$v(x, y, z) = \begin{bmatrix} \alpha \cdot I(x, y, z) \\ \beta \cdot x \\ \beta \cdot y \\ \beta \cdot z \end{bmatrix}$$

$$\rightarrow D(v_i, v_j) = \|v_i - v_j\|_1$$

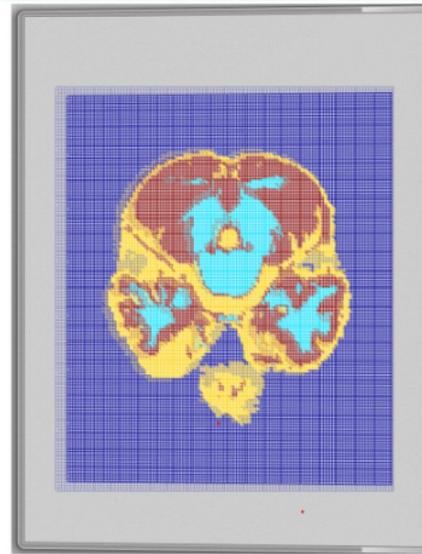
$$D(v_i, v_j) = \|v_i - v_j\|_2$$

- Means of K-means: Multiple layer of K-means



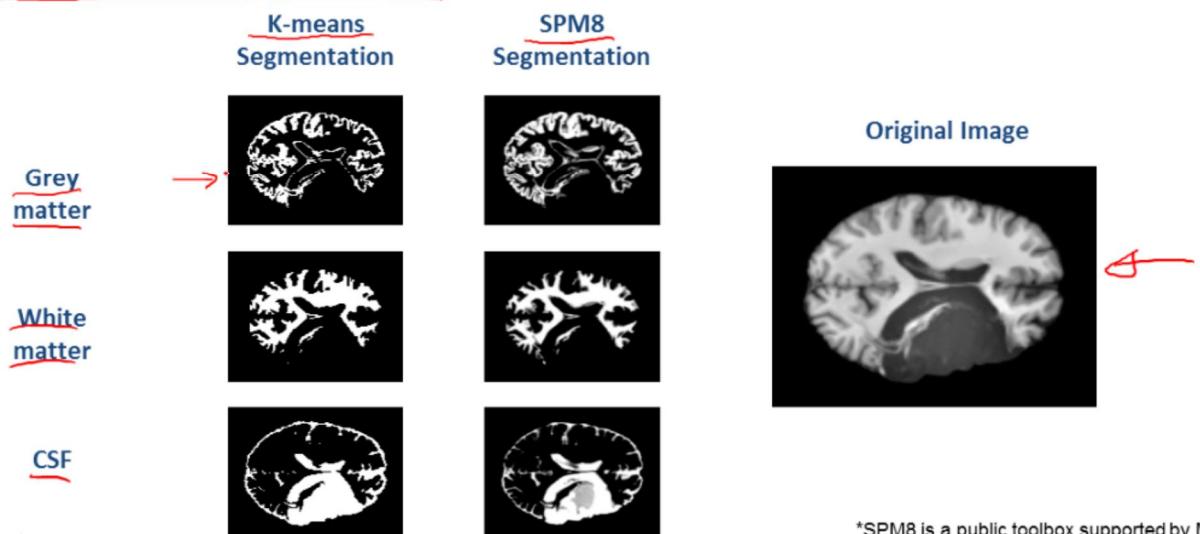
The K-means Clustering

3D segmentation of
A human brain

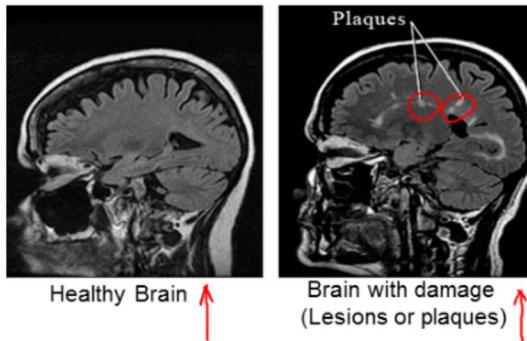


Experimental Results

3D segmentation by K-means on slice 60, compared to SPM8* segmentation result



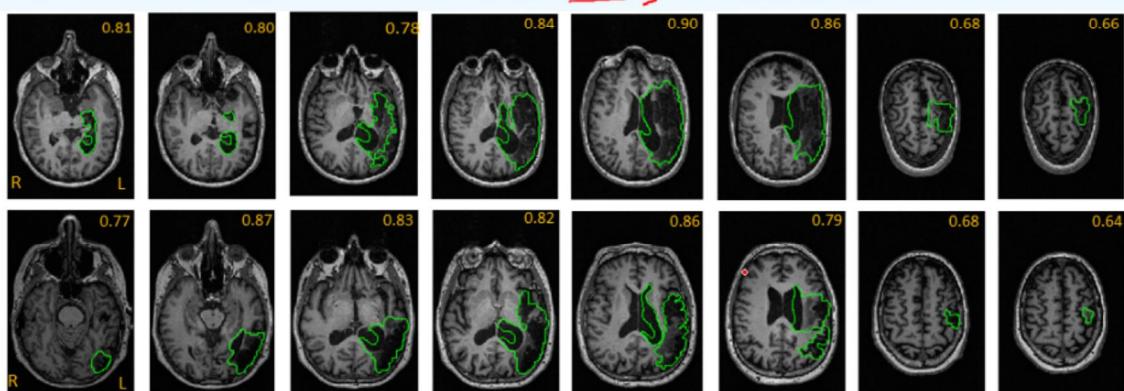
**Brain Image Segmentation is a very difficult problem;
the identification of lesions is even harder**



Current Research Status:

- Trying to identify the stroke lesion and the perilesional region to predict recovery or response to therapy.
- Typically this is drawn by hand which is time consuming and subjective process.

Segmentation Techniques may Save Lives



- Used proactively, the technique can screen a healthy individual's brain to detect abnormalities.
- Following a disease, however, the same **technique** can assess the patient's response to therapy and predict recovery

$$\text{Dice Similarity Index} \quad s = \frac{2|X \cap Y|}{|X| + |Y|}$$

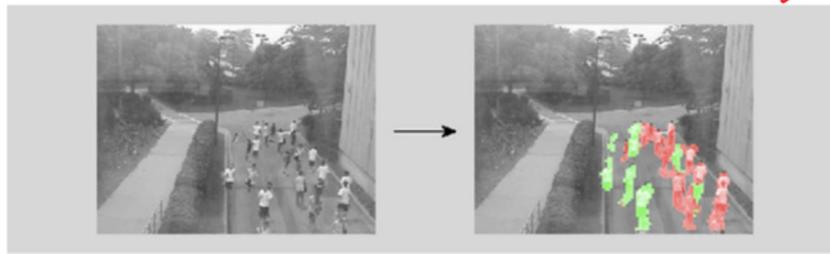
$[0, 1]$

Motion-based Segmentation

Motion is a powerful cue used by humans and many other animals to extract objects or regions of interest from a background of irrelevant detail.

Given an image sequence, determine

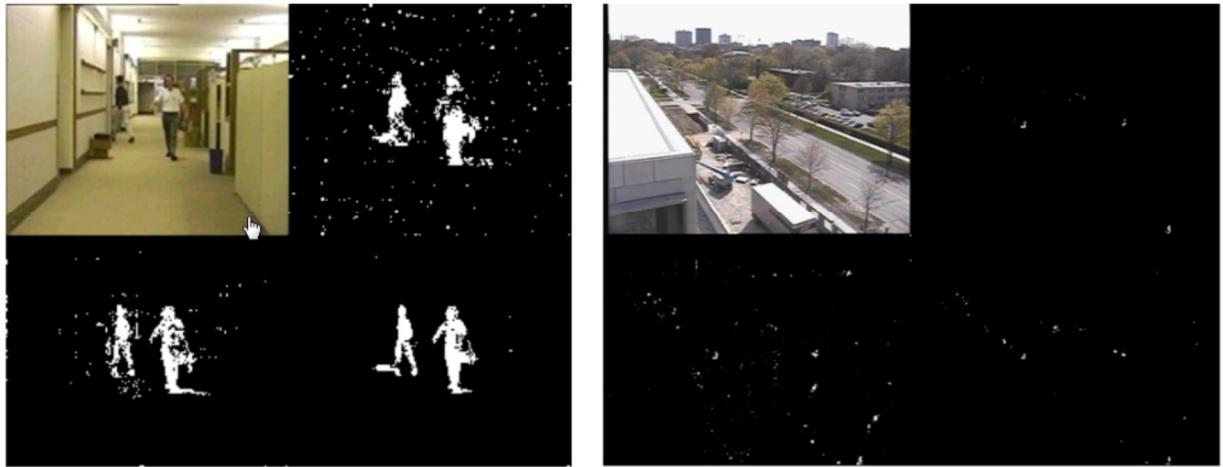
- **Number** of motion models (affine, Euclidean, etc.)
- **Motion model**: affine (2D) or Euclidean (3D)
- **Segmentation**: model to which each pixel belongs



Motion-based Segmentation

- Object detection
 - pedestrian detection
- Tracking
 - vehicle tracking
- Robotics
- Surveillance
- Image and video compression
- Scene reconstruction
- Video manipulation / editing
 - video matting
 - video annotation
 - motion magnification

Example

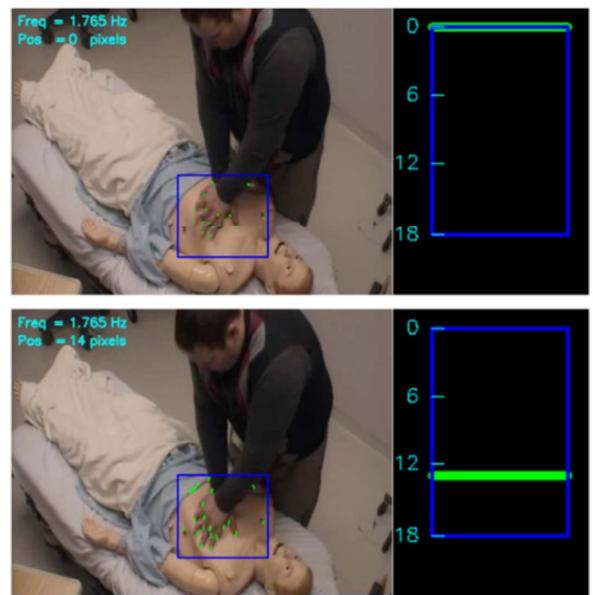


Motion-based Segmentation

Segmenting images based on **common motion**

Points moving together are **grouped together**

Other useful information can be extracted through the process,
e.g. depth and frequency of motion



Motion-based Segmentation

Basic Approach

- A difference image between t_i and t_j is defined as

$$d_{i,j}(x, y) = \begin{cases} 1 & \text{if } |f(x, y, i) - f(x, y, j)| > T \\ 0 & \text{otherwise} \end{cases}$$

T is a specified threshold.

- Object motions result in 1's in $d_{i,j}$

- Assumptions
 - Images registered spatially
 - Relatively constant Illumination

Motion-based Segmentation

Improvement by Accumulative difference image (ADI)

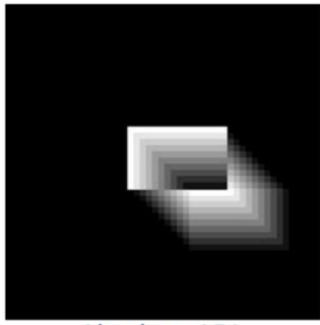
$$A_k(x, y) = \begin{cases} A_{k-1}(x, y) + 1 & \text{if } |R(x, y) - f(x, y, k)| > T \\ A_{k-1}(x, y) & \text{otherwise} \end{cases} \quad \text{Absolute}$$

$$P_k(x, y) = \begin{cases} P_{k-1}(x, y) + 1 & \text{if } [R(x, y) - f(x, y, k)] > T \\ P_{k-1}(x, y) & \text{otherwise} \end{cases} \quad \text{Positive}$$

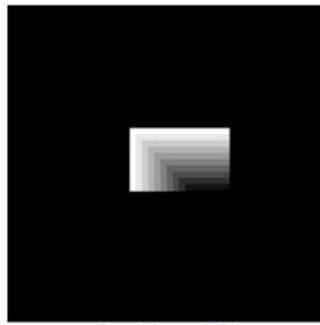
$$N_k(x, y) = \begin{cases} N_{k-1}(x, y) + 1 & \text{if } [R(x, y) - f(x, y, k)] < -T \\ N_{k-1}(x, y) & \text{otherwise} \end{cases} \quad \text{Negative}$$

- Isolated entries in $d_{i,j}$ resulting from noise are not insignificant
- Although thresholding using connectivity may remove these entries, small or slow-moving objects are also removed.
- Use of ADI ignores changes that occur only sporadically

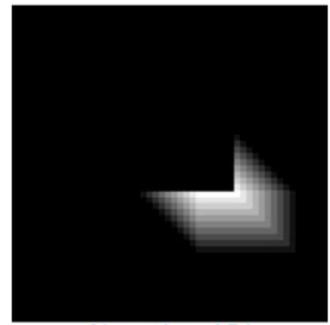
Motion-based Segmentation



Absolute ADI



Positive ADI



Negative ADI

From positive ADI : size of the moving object

location of the moving object in reference image

From absolute ADI / negative ADI : speed and direction of the motion

Mean Shift

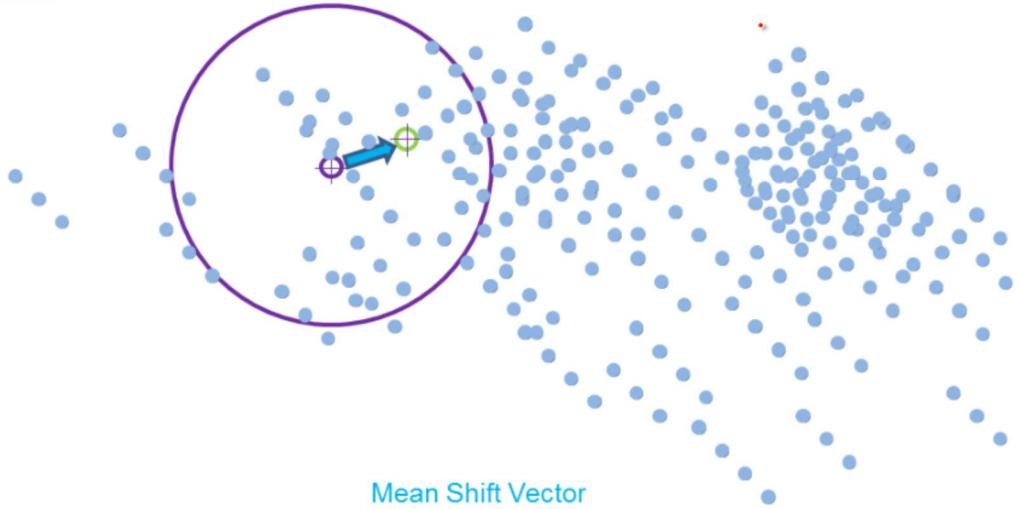
-
- A scatter plot of blue dots representing data points. A red circle highlights a cluster of points, with a red arrow pointing towards it, indicating the center of mass or mean shift.
- Data transferred into feature space
 - Feature space : color, texture, gradient etc.
 - Find modes/ local maximum of density in the feature space

Mean Shift

Region of Interest

Center of Mass

Mean Shift Vector



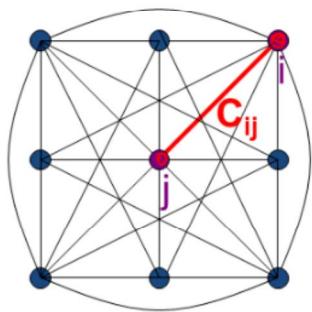
Mean Shift



Mean Shift



Graph Cut



Images as graphs

Pixels correspond to Vertices (V)

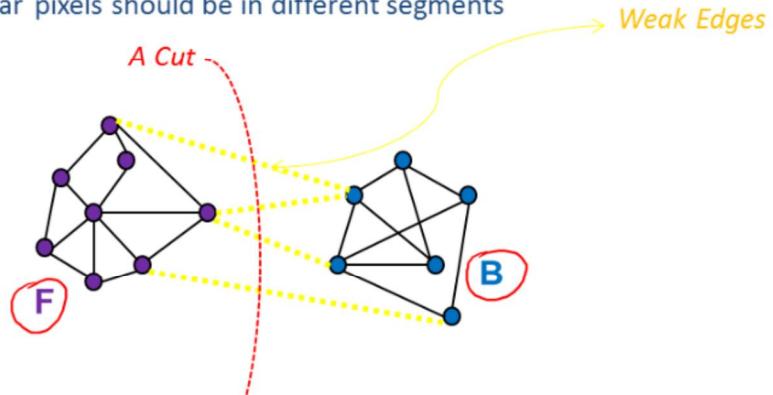
Pairs of pixels correspond to Edges (E)

Similarities between pixels correspond to weights/costs $C_{i,j}$

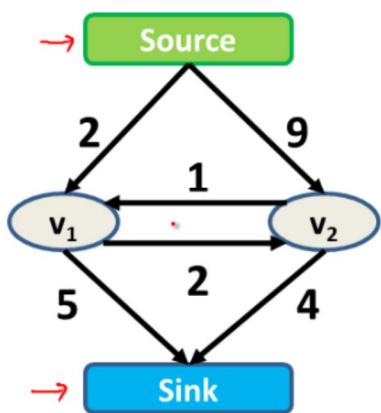
Graph Cut

Intuition : Break Graph into Segments

- Delete Edges that cross between segments
- Easiest to break Edges that have low cost (low similarity)
 - similar pixels should be in the same segments
 - dissimilar pixels should be in different segments

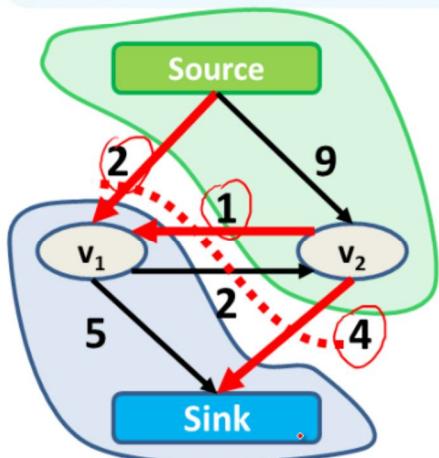


Example: st-Mincut Problem



What is a st-cut?

st-Mincut Problem



$$2 + 1 + 4 = 7$$

What is a st-cut?

An st-cut (S, T) divides the nodes between source and sink.

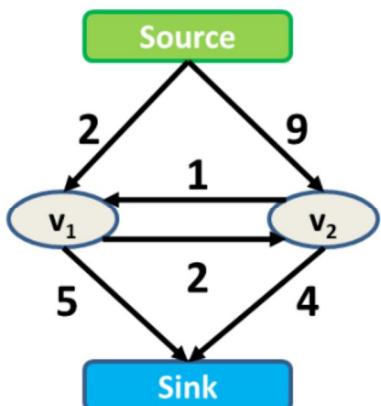
What is the cost of a st-cut?

Sum of cost of all edges going from S to T

What is the st-mincut?

st-cut with the minimum cost

How to compute the st-mincut?



Solve the dual maximum flow problem

Compute the maximum flow between Source and Sink

Constraints

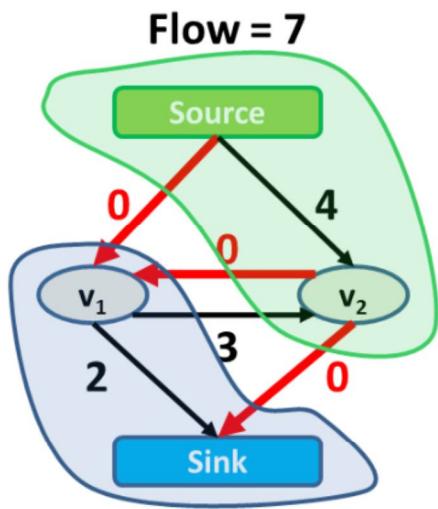
Edges: Flow < Capacity

Nodes: Flow in = Flow out

Min-cut\Max-flow Theorem

In every network, the maximum flow equals the cost of the st-mincut

Maxflow Algorithms



Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity*
2. Push maximum possible flow through this path
3. Repeat until no path can be found

*Algorithms assume non-negative capacity

Example: using Min Cut

